



Exam written 24012012 Solutions

Pattern Recognition (Technische Universiteit Delft)

IN4085
Pattern Recognition

Written examination
24-01-2012, 14:00-17:00

- There are five questions
- You have 30 minutes to answer the first question (Answer sheets 1-2), during which you cannot consult the book or any other material
- After you have handed in Answer sheets 1-2, you are allowed to consult any material you have brought to answer the following questions in the remaining 150 minutes
- Answer **EACH QUESTION** on a **SEPARATE SHEET OF PAPER**
- As much as possible, include the **CALCULATIONS** you made to get to an answer
- Do not forget to put your **NAME** and **STUDENT NUMBER** on top of every sheet
- Do not forget to hand this exam, including all Answer sheets

Answer sheet 1

Name :
Student number :

1 Statements

(10 points)

Circle the correct statement, i.e. TRUE or FALSE. If a statement does not hold in general, but only under certain conditions that are not mentioned, then the statement should be marked as FALSE. 10 correct answers will give you 0 points; each additional correct answer gives you 1 point.

Classification

1. To obtain a classifier with a high classification performance, you need to estimate the class-conditional probabilities well.
TRUE FALSE
Solution: FALSE
2. For a given dataset, the Bayes' classifier has the lowest classification error.
TRUE FALSE
Solution: TRUE
3. The linear discriminant analysis `ldc` is scale insensitive.
TRUE FALSE
Solution: TRUE
4. A confusion matrix shows how expensive it is to misclassify an object from one class to another class.
TRUE FALSE
Solution: FALSE
5. The logistic classifier minimizes the number of erroneously classified objects of the training set.
TRUE FALSE
Solution: FALSE
6. The decision boundary of a two-class classifier is defined as all points \mathbf{x} for which holds:
 $p(\mathbf{x}|\omega_1) = p(\mathbf{x}|\omega_2)$.
TRUE FALSE
Solution: FALSE
7. The Receiver Operating Characteristic Curve is independent of class priors in the test set.
TRUE FALSE
Solution: TRUE
8. When the classes in a classification problem have a Gaussian distribution, the optimal classifier is the quadratic classifier `qdc`.
TRUE FALSE
Solution: FALSE
9. When k in the k -nearest neighbor classifier is set to the number of training examples, $k = n$, the classifier becomes independent of class distributions.
TRUE FALSE
Solution: TRUE

10. Increasing the number of folds in crossvalidation reduces the variance in the error estimate.
TRUE FALSE
Solution: TRUE

Answer sheet 2

Name :
Student number :

Clustering

1. Hierarchical clustering makes no assumptions about the underlying distribution of the data.
TRUE FALSE
Solution: TRUE
2. A hierarchical clustering result is sensitive to its random initialization.
TRUE FALSE
Solution: FALSE
3. The larger K is in K -means clustering the better the model typically fits to the training data.
TRUE FALSE
Solution: TRUE
4. If single linkage clustering fails to work on a particular data set, complete linkage will work (and vice versa).
TRUE FALSE
Solution: FALSE
5. To avoid local minima, the K -means algorithm should always be started with the K means centered on randomly selected objects.
TRUE FALSE
Solution: FALSE

Feature Extraction and Selection

1. Reducing the feature vector dimensionality by means of principle component analysis (PCA) will lead to worse classification performance because it is an unsupervised technique.
TRUE FALSE
Solution: FALSE
2. Extracting features can never lead to a decrease of the Bayes error.
TRUE FALSE
Solution: TRUE
3. In feature selection, the number of different feature sets of size 7 that one can select from 1,000 dimensions is larger than the number of different feature sets of size 993 that one can select.
TRUE FALSE
Solution: FALSE
4. PCA assumes that the underlying data distribution is Gaussian.
TRUE FALSE
Solution: FALSE
5. Even by extracting more than 5 (linear or nonlinear) features from a 5-dimensional classification problem, one cannot reduce the Bayes error of the original problem.
TRUE FALSE
Solution: TRUE

2 Classification: probabilities

(10 points)

A classifier obtained the following class-conditional probabilities for ten test objects:

object nr	Class conditional probabilities		true class
	class 1	class 2	
1	0.07	0.02	1
2	0.05	0.04	1
3	0.1	0.2	1
4	0.3	0.35	1
5	0.8	0.07	1
6	0.02	0.01	1
7	0.2	0.5	2
8	0.4	0.6	2
9	0.03	0.8	2
10	0.01	0.2	2

- a. Assume equal class priors. How many objects are misclassified? (1 points)

Solution: 2.

- b. Assume equal class priors. What is the classification error of this test set? (2 points)

Solution: $\varepsilon = 0.5 \cdot 2/6 + 0.5 \cdot 0 = 0.167$.

- c. Assume that class 1 is ten times as likely as class 2, $p(\omega_1) = 10p(\omega_2)$. What is the classification error now? (3 points)

Solution: $\varepsilon = 0.909 \cdot 0 + 0.091 \cdot 2/4 = 0.046$.

- d. Assume still the case that $p(\omega_1) = 10p(\omega_2)$. Will the rejection of the most uncertain object improve the performance of the classifier? (Assume that the cost of rejection is lower than the cost of misclassification) (2 points)

Solution: No.

- e. Another classifier obtained the following class-conditional probabilities for the ten objects:

object nr	Class conditional probabilities		true class
	class 1	class 2	
1	0	0	1
2	0	0	1
3	0	0	1
4	0.0005	0	1
5	0	0	1
6	0	0	1
7	0	0	2
8	0	0	2
9	0	0	2
10	0	0	2

Explain what happened here. (2 points)

Solution: It seems that the classifier overfitted on the training objects. The test objects are far from (or at least, not identical to) the training objects, so the class conditional probabilities are all zero. This classifier does not generalize at all.

3 Classification: curves

(10 points)

- a. Assume we train a Parzen classifier on some training data, and we vary the width parameter h . First make a plot of the likelihood of the training data as a function of the width parameter h . Second, explain how you can optimize the width parameter h . (1 points)

Solution: The smaller you make h , the higher the likelihood. On the training set it suggests that $h = 0$ is optimal, but of course, you need an independent test set to get a good idea. Then you probably see for a $h > 0$ an optimum.

- b. Draw a learning curve. Show both the true error as the apparent error. Name the axes, and list the important characteristics of the curve. Mention also the words "Bayes' error" and "overfitting". (2 points)

Solution: 0. Error vs. training set size. 1. Higher true error than apparent error. 2. Errors converge to asymptotic error. 3. Difference between true and apparent error is overfitting. 4. Bayes' error is lower than the asymptotic error.

- c. Assume we have a training set consisting of n objects and d features with 2 classes. How many parameters have to be estimated for the quadratic classifier? And how many parameters have to be estimated for the linear support vector classifier? And how many for the Parzen classifier? (2 points)

Solution: For the qdc: 1 for the class prior, $2d$ for the means, and two covariance matrices: $2 \cdot \frac{1}{2}d(d+1)$, so in total $N_{qdc} = 1 + 2d + d(d+1) = d^2 + 3d + 1$. For the SVM we need to estimate the α 's, so that is $N_{svm} = n$. But you may also consider the optimization of C , so then you get $N_{svm} = n + 1$. For the Parzen classifier we only have to optimize the width parameter h .

- d. Assume we have a two-class classification problem, with a fixed number of n training points, but with a very large set of weakly informative features. Assume we randomly select features from the large feature set. Make a plot of the classification error as a function of the number of features, for the quadratic classifier, the linear support vector classifier and the Parzen classifier. Draw them all in one plot, and explain the similarities and differences between the curves. (3 points)

Solution: I expect the support vector to steadily decrease, but it may not get the overall best performance (because it is only linear). The Parzen classifier may decrease a lot in the beginning, probably getting the lowest error, but will also collapse for larger feature sizes. The quadratic will probably be a bit worse than the Parzen, and will also collapse when the dimensionality reaches the number of training points per class.

- e. Assume we are in the same situation as in previous question d. and we decide to combine the three classifiers using *majority voting*. Again plot the error as a function of the number of features, for all classifiers and the combined classifier. (2 points)

Solution: You hope that in the beginning, all classifiers contribute something independent, and that combining results in better performance. At the point that the qdc crashes, you get a majority vote between Parzen and svm, probably resulting in a performance somewhere in between?? When both Parzen and qdc crash, also the combining does not work anymore.

4 Feature Extraction and Selection

(11 points)




Figure 3: Example configuration of three classes.

Let us consider a two-dimensional, three-class problem in which all classes are uniformly distributed in a disc with radius 1. All class means are variable but differ so that there is no class overlap in the 2D feature space. The prior probabilities of all classes are assumed equal. Figure 3 above gives an example configuration in which the three class means are given by $(2, 3)$, $(5, 7)$, and $(10, 6)$, respectively.

- Assume we reduce the feature dimensionality from 2 to 1 by means of the Fisher mapping (**fisherm**). Configure the three class means such that in the 2D space there is *no overlap* but in the 1D space found by the Fisher mapping two of the three classes *completely* overlap. You are only allowed to put the class means in the *grid points* given in the figure. Provide the three coordinates of your solution. (3 points)

Solution: There are many possible solutions, though the “means on grid” restriction limits the possibilities drastically. One solution is given by $(3, 1)$, $(1, 3)$, and the third disc somewhere on $(5, 5)$, or $(6, 6)$, etc.

- Is it possible to give three means for which there is no overlap in 2D but for which all three classes overlap completely in 1D when using the Fisher mapping? (2 points)

Solution: No, this is impossible. For them to overlap perfectly, they have to be perfectly aligned, but this means that exactly the wrong, i.e., the most optimal direction, will be chosen by Fisher.




Figure 4: Example configuration of two classes (both of which in turn consist of two clusters).

We now consider a similar setting as exemplified in Figure 3, but now we have *two* classes (one black one white) both consisting of two clusters. All four clusters are uniformly distributed in a disc with radius 1. (The four clusters have equal priors.) An example configuration is given in Figure 4.

- c. Configure the four class means such that if one would look at either of the two features, both classes would perfectly *overlap*, while in the original 2D space the two classes are perfectly non-overlapping. In other words, construct an example for which selecting a single feature would give a Bayes error of 0.5, while the Bayes error in the original space is 0.

Again, you are only allowed to put the means in the *grid points* given in the figure. Provide the four coordinates of the clusters of your solution. Make sure you make clear where the black clusters go and where the white go. (2 points)

Solution: There are many possible solutions, though the “means on grid” restriction limits the possibilities. The only thing one needs to do is pick any rectangle aligned with the axes and put the cluster means in its vertices making sure that clusters from one class are in opposing vertices.

- d. Would a linear feature extraction technique be able to provide a better one-dimensional subspace for the problem created in a. than feature selection is able to do? That is, can feature extraction for a 1D feature for which the two classes do not fully overlap? (2 points)

Solution: Yes.

- e. Is it possible to create a classification problem, again positioning the four clusters, such that feature selection will outperform any kind of feature extraction? Explain your answer. (2 points)

Solution: “Obviously” not.

5 Clustering

(9 points)



Figure 5: A scatter plot of three hundred feature vectors in a two-dimensional feature space.

Figure 5 above displays three hundred points in 2D drawn from three normal distributions. We want to cluster the points in *two* clusters by means of K -means clustering with $K = 2$, i.e., we want to divide the data set up in two groups.

- a. If we initialize K -means clustering with the two crosses indicated in the same figure, how will the final clustering look like (indicate this clustering in figure 6 or describe it in some other unambiguous way)? (2 points)

Solution: The two nearest clusters form one cluster, while the distant cluster will remain on its own.

Now consider all possible random initializations for K -means clustering (with K still equal to 2) for the problem in the figure.

- b. How many different clusterings consisting of *two nonempty* clusters can be found when randomly initializing the two initial cluster centers? (3 points)

Solution: One only. There is no other (stable?) solutions possible then the one from **a**.

- c. Will every random initialization, with two initial means, result in a solution with two clusters? Explain your answer. (2 points)

Solution: No, one could also end up with a single cluster.

- d. If we choose to take four cluster centers ($K = 4$), how can one detect that this is probably a wrong number of clusters? What measure or test reveals this? (2 points)

Solution: Rerunning K -means a couple of times and checking the cluster stability is one option.

Answer sheet 3

Name :
Student number :



Figure 6: A scatter plot of three hundred feature vectors in a two-dimensional feature space.